

Project Name - HR Analytics Project

Name - Aman Mulla.

Batch - DS2307

Project Summary -

The proposed project centers on HR Analytics' role in addressing employee attrition and enhancing organizational effectiveness. Employee turnover imposes substantial costs on companies, including expenses related to hiring, training, and a decline in cumulative knowledge. Furthermore, high attrition hampers the establishment of a stable, experienced workforce, potentially impacting customer relations and operational efficiency. HR Analytics offers a strategic approach by leveraging data insights to identify attrition patterns, understand contributing factors, and implement proactive measures for retention. By analyzing employee data, such as performance metrics, engagement levels, and workplace dynamics, HR Analytics facilitates informed decision-making to optimize retention strategies. The project aims to explore these dynamics, illustrating how data-driven insights from HR Analytics can mitigate attrition's adverse impacts while fostering a more robust, engaged workforce.

HR Analytics:

Human resource analytics (HR analytics) is an area in the field of analytics that refers to applying analytic processes to the human resource department of an organization in the hope of improving employee performance and therefore getting a better return on investment. HR analytics does not just deal with gathering data on employee efficiency. Instead, it aims to provide insight into each process by gathering data and then using it to make relevant decisions about how to improve these processes.

Attrition in HR:

Attrition in human resources refers to the gradual loss of employees overtime. In general, relatively high attrition is problematic for companies. HR professionals often assume a leadership role in designing company compensation programs, work culture, and motivation systems that help the organization retain top employees.

We have below columns in our dataset,

Sure, here are brief descriptions of each input feature:

1. **Age:** Employee's age.
2. **Attrition:** Indicates whether the employee has left the company ('Yes' or 'No').
3. **BusinessTravel:** Frequency of travel for work ('Travel_Rarely', 'Travel_Frequently', 'Non-Travel').
4. **DailyRate:** Employee's daily rate of pay.
5. **Department:** Department in the company where the employee works ('Sales', 'Research & Development', 'Human Resources').
6. **DistanceFromHome:** Distance between employee's home and workplace.
7. **Education:** Employee's level of education (from 1 to 5).
8. **EducationField:** Field of education the employee studied.
9. **EmployeeCount:** Number of employees, typically constant (might not provide useful information).
10. **EmployeeNumber:** Unique identifier for each employee.
11. **EnvironmentSatisfaction:** Employee's satisfaction with the work environment (from 1 to 4).
12. **Gender:** Employee's gender ('Male' or 'Female').
13. **HourlyRate:** Employee's hourly rate of pay.
14. **JobInvolvement:** Level of job involvement (from 1 to 4).
15. **JobLevel:** Employee's job level (from 1 to 5).
16. **JobRole:** Role or position in the company.
17. **JobSatisfaction:** Employee's satisfaction with their job (from 1 to 4).
18. **MaritalStatus:** Employee's marital status ('Single', 'Married', 'Divorced').
19. **MonthlyIncome:** Employee's monthly income.
20. **MonthlyRate:** Employee's monthly rate of pay.
21. **NumCompaniesWorked:** Number of companies the employee has worked for.
22. **Over18:** Indicates if the employee is over 18 years old ('Yes').
23. **OverTime:** Indicates if the employee works overtime ('Yes' or 'No').
24. **PercentSalaryHike:** Percentage of the employee's salary increase.
25. **PerformanceRating:** Employee's performance rating (from 3 to 4).
26. **RelationshipSatisfaction:** Employee's satisfaction with work relationships (from 1 to 4).
27. **StandardHours:** Standard number of working hours (might not provide useful information).
28. **StockOptionLevel:** Level of stock options the employee has (from 0 to 3).
29. **TotalWorkingYears:** Total number of years the employee has worked.
30. **TrainingTimesLastYear:** Number of training times attended by the employee last year.

- 31. **WorkLifeBalance**: Employee's satisfaction with work-life balance (from 1 to 4).
- 32. **YearsAtCompany**: Number of years the employee has worked at the company.
- 33. **YearsInCurrentRole**: Number of years the employee has been in the current role.
- 34. **YearsSinceLastPromotion**: Number of years since the employee's last promotion.
- 35. **YearsWithCurrManager**: Number of years the employee has been with the current manager.

▼ **Problem Statement**

High employee attrition rates pose significant challenges for organizations, resulting in increased operational costs, a loss of cumulative expertise, and potential disruptions in customer service. This project aims to investigate the impact of attrition on companies and explore how HR Analytics can effectively analyze and address this issue. The objective is to understand the factors contributing to employee turnover, identify patterns or trends within the data, and develop proactive strategies to minimize attrition. By leveraging data-driven insights, this study seeks to offer actionable recommendations for HR professionals and organizational leaders to optimize retention efforts and enhance overall workforce stability and performance.

▼ **Knowing data and variable in dataset**

```
# Importing Necessary Libraries

import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt

pd.set_option('display.max_rows', None)
pd.set_option('display.max_columns', None)

HR_data = pd.read_csv('/content/drive/MyDrive/DataSets/WA_Fn-UseC_-HR-Employee-Attrition.csv')

HR_data.head()
```

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	EducationField	EmployeeCount	EmployeeNumbe
0	41	Yes	Travel_Rarely	1102	Sales	1	2	Life Sciences	1	
1	49	No	Travel_Frequently	279	Research & Development	8	1	Life Sciences	1	
2	37	Yes	Travel_Rarely	1373	Research & Development	2	2	Other	1	
3	33	No	Travel_Frequently	1392	Research & Development	3	4	Life Sciences	1	
4	27	No	Travel_Rarely	591	Research & Development	2	1	Medical	1	

```
HR_data.columns

Index(['Age', 'Attrition', 'BusinessTravel', 'DailyRate', 'Department',
      'DistanceFromHome', 'Education', 'EducationField', 'EmployeeCount',
      'EmployeeNumber', 'EnvironmentSatisfaction', 'Gender', 'HourlyRate',
      'JobInvolvement', 'JobLevel', 'JobRole', 'JobSatisfaction',
      'MaritalStatus', 'MonthlyIncome', 'MonthlyRate', 'NumCompaniesWorked',
      'Over18', 'OverTime', 'PercentSalaryHike', 'PerformanceRating',
      'RelationshipSatisfaction', 'StandardHours', 'StockOptionLevel',
      'TotalWorkingYears', 'TrainingTimesLastYear', 'WorkLifeBalance',
      'YearsAtCompany', 'YearsInCurrentRole', 'YearsSinceLastPromotion',
      'YearsWithCurrManager'],
      dtype='object')

# Will Check for shape of dataset

HR_data.shape

(1470, 35)
```

We have total 1470 rows and 35 column in our detaset.

Dataset Information

```
HR_data.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1470 entries, 0 to 1469
Data columns (total 35 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Age                                    1470 non-null   int64
1   Attrition                            1470 non-null   object
2   BusinessTravel                        1470 non-null   object
3   DailyRate                            1470 non-null   int64
4   Department                            1470 non-null   object
5   DistanceFromHome                     1470 non-null   int64
6   Education                             1470 non-null   int64
7   EducationField                        1470 non-null   object
8   EmployeeCount                         1470 non-null   int64
9   EmployeeNumber                       1470 non-null   int64
10  EnvironmentSatisfaction               1470 non-null   int64
11  Gender                               1470 non-null   object
12  HourlyRate                           1470 non-null   int64
13  JobInvolvement                       1470 non-null   int64
14  JobLevel                             1470 non-null   int64
15  JobRole                              1470 non-null   object
16  JobSatisfaction                       1470 non-null   int64
17  MaritalStatus                        1470 non-null   object
18  MonthlyIncome                        1470 non-null   int64
19  MonthlyRate                          1470 non-null   int64
20  NumCompaniesWorked                   1470 non-null   int64
21  Over18                              1470 non-null   object
22  OverTime                             1470 non-null   object
23  PercentSalaryHike                    1470 non-null   int64
24  PerformanceRating                    1470 non-null   int64
25  RelationshipSatisfaction              1470 non-null   int64
26  StandardHours                        1470 non-null   int64
27  StockOptionLevel                     1470 non-null   int64
28  TotalWorkingYears                    1470 non-null   int64
29  TrainingTimesLastYear                1470 non-null   int64
30  WorkLifeBalance                      1470 non-null   int64
31  YearsAtCompany                       1470 non-null   int64
32  YearsInCurrentRole                   1470 non-null   int64
33  YearsSinceLastPromotion               1470 non-null   int64
34  YearsWithCurrManager                 1470 non-null   int64
dtypes: int64(26), object(9)
memory usage: 402.1+ KB
```

From .info(), we can observe that there were variables with datatype of float,object and int only.

```
# Will check for description of dataset

HR_data.describe()
```

	Age	DailyRate	DistanceFromHome	Education	EmployeeCount	EmployeeNumber	EnvironmentSatisfaction	HourlyRate	...
count	1470.000000	1470.000000	1470.000000	1470.000000	1470.0	1470.000000	1470.000000	1470.000000	
mean	36.923810	802.485714	9.192517	2.912925	1.0	1024.865306	2.721769	65.891156	
std	9.135373	403.509100	8.106864	1.024165	0.0	602.024335	1.093082	20.329428	
min	18.000000	102.000000	1.000000	1.000000	1.0	1.000000	1.000000	30.000000	
25%	30.000000	465.000000	2.000000	2.000000	1.0	491.250000	2.000000	48.000000	
50%	36.000000	802.000000	7.000000	3.000000	1.0	1020.500000	3.000000	66.000000	
75%	43.000000	1157.000000	14.000000	4.000000	1.0	1555.750000	4.000000	83.750000	
max	60.000000	1499.000000	29.000000	5.000000	1.0	2068.000000	4.000000	100.000000	

From .describe() we can get count, mean, minimum value, maximum values and quartile value for each numerical column.

```
HR_data.isnull().sum()

Age                                0
Attrition                          0
BusinessTravel                     0
DailyRate                          0
Department                          0
DistanceFromHome                   0
Education                          0
EducationField                     0
EmployeeCount                       0
EmployeeNumber                     0
```

```

EnvironmentSatisfaction    0
Gender                    0
HourlyRate                0
JobInvolvement            0
JobLevel                  0
JobRole                   0
JobSatisfaction            0
MaritalStatus             0
MonthlyIncome             0
MonthlyRate               0
NumCompaniesWorked        0
Over18                    0
OverTime                  0
PercentSalaryHike         0
PerformanceRating         0
RelationshipSatisfaction  0
StandardHours             0
StockOptionLevel          0
TotalWorkingYears         0
TrainingTimesLastYear     0
WorkLifeBalance           0
YearsAtCompany            0
YearsInCurrentRole        0
YearsSinceLastPromotion   0
YearsWithCurrManager      0
dtype: int64

```

```
sns.heatmap(HR_data.isnull())
```




Chart - 1

Attrition Distribution

```
attrition_distribution = HR_data['Attrition'].value_counts()
```

```
attrition_distribution
```

```
labels = attrition_distribution.index
```


```
plt.figure(figsize=(8,6))
```

```
sns.set_palette('pastel')
```

```
plt.pie(attrition_distribution, labels=attrition_distribution.index, autopct='%0f%%')
```

```
plt.legend(labels, title="Type", loc="upper right", bbox_to_anchor=(1, 0, 0.5, 1))
```

```
plt.show()
```

**Insights from above chart:**


- Displays the distribution of attrition types using a pie chart.
- The legend showing the labels (types of attrition - likely 'Yes' and 'No').
- This visualization includes the proportion or percentage of employees who have experienced attrition ('Yes') compared to those who have not ('No'). This helps in understanding the attrition rate within the dataset, providing a quick view of how many employees have left the company versus those who have stayed.

✓ Chart - 2

Age Distribution

```
sns.violinplot(x=HR_data["Age"],linewidth=1, linecolor="k")
```

```
plt.xlabel('Age')  
plt.ylabel('Count of Employees')  
plt.title('Age Distribution')  
plt.show()
```

**Insights from above plot:**

- The plot will show a distribution of ages among employees in your dataset. The spread and concentration of ages can be visualized through the shape of the violins.
- The width of the plot at different points indicates the density of employees in specific age groups. A wider section suggests more employees falling within that age range.
- A symmetric plot implies a more uniform distribution, while skewness can indicate an uneven distribution toward older or younger ages.

Chart - 3

Department-wise Attrition

```
department_attrition_count = HR_data.groupby(['Department', 'Attrition']).size().unstack()
```

```
department_attrition_count.plot(kind='bar', stacked=True, figsize=(10, 6))
```

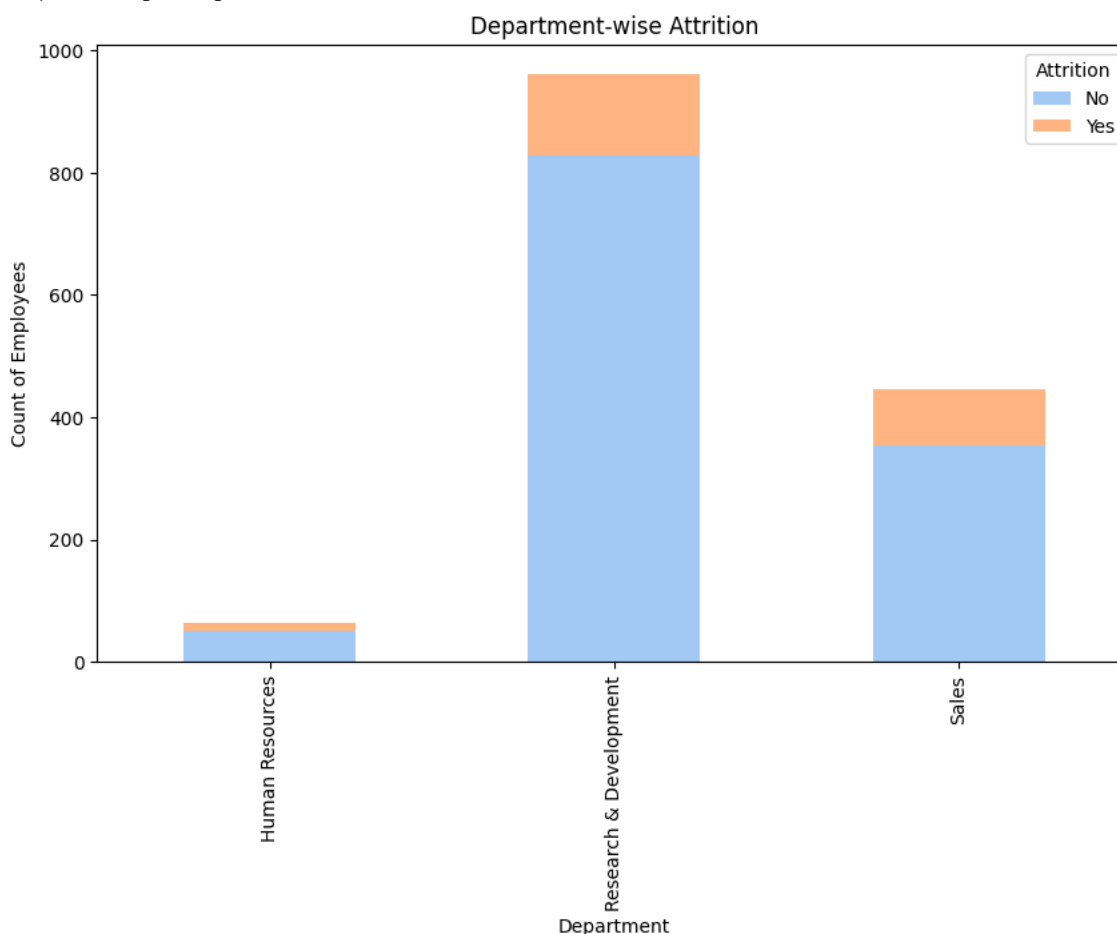
```
plt.xlabel('Department')
```

```
plt.ylabel('Count of Employees')
```

```
plt.title('Department-wise Attrition')
```

```
plt.legend(title='Attrition')
```

<matplotlib.legend.Legend at 0x7f34830d97b0>

**Insights from above plot:**

- Each department will have segments representing both "Attrition" and "No Attrition" categories.
- By comparing the length of the bars for each department, gauge the overall size of the workforce in each department. Larger bars indicate a higher count of employees in that specific department.
- Within each department bar, the stacked segments will show the proportion of employees who have experienced attrition versus those who have not.

Chart - 4

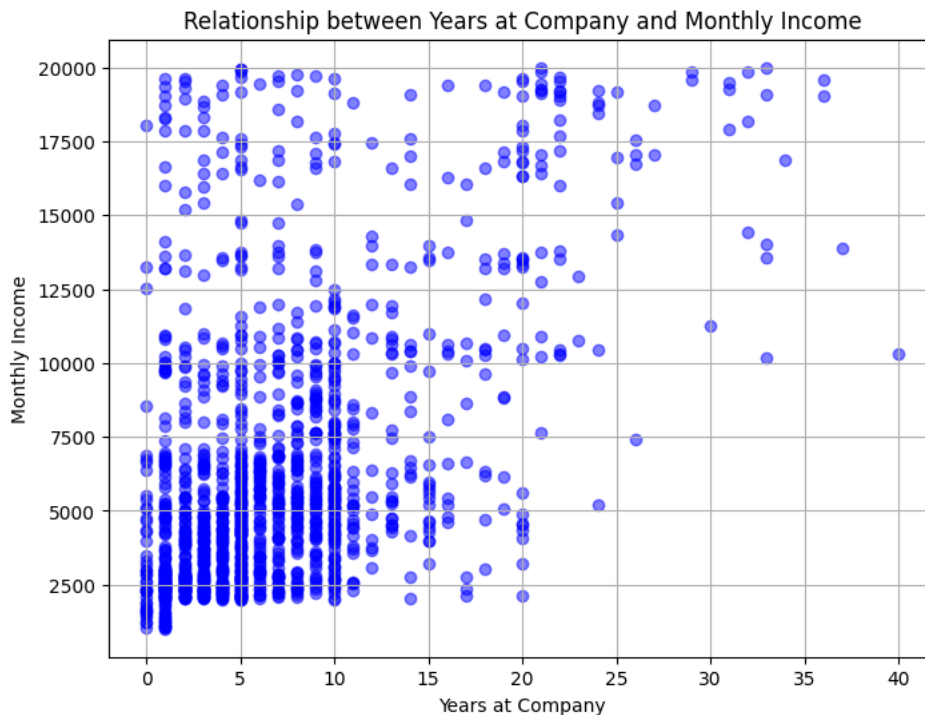
Years at Company vs. Salary

```
# Extracting relevant columns
years_at_company = HR_data['YearsAtCompany']
monthly_income = HR_data['MonthlyIncome']

# Creating the scatter plot
plt.figure(figsize=(8, 6))
plt.scatter(years_at_company, monthly_income, color='blue', alpha=0.5)

# Adding labels and title
plt.xlabel('Years at Company')
plt.ylabel('Monthly Income')
plt.title('Relationship between Years at Company and Monthly Income')

# Show plot
plt.grid(True)
plt.show()
```



Insights from above plot:

- Upon observing the scatter plot, you may notice that there isn't a clear, direct linear relationship between "Years at Company" and "Monthly Income."
- There might be clusters or patterns visible in the data points. For example, you might notice groups of employees who have similar monthly incomes despite varying years of experience at the company. This could indicate the presence of categorical factors like job roles, departments, or education levels that impact income more significantly.
- The absence of a noticeable trend in the scatter plot might indicate that the relationship between these two variables is more complex or influenced by additional factors.

✓ Chart - 5

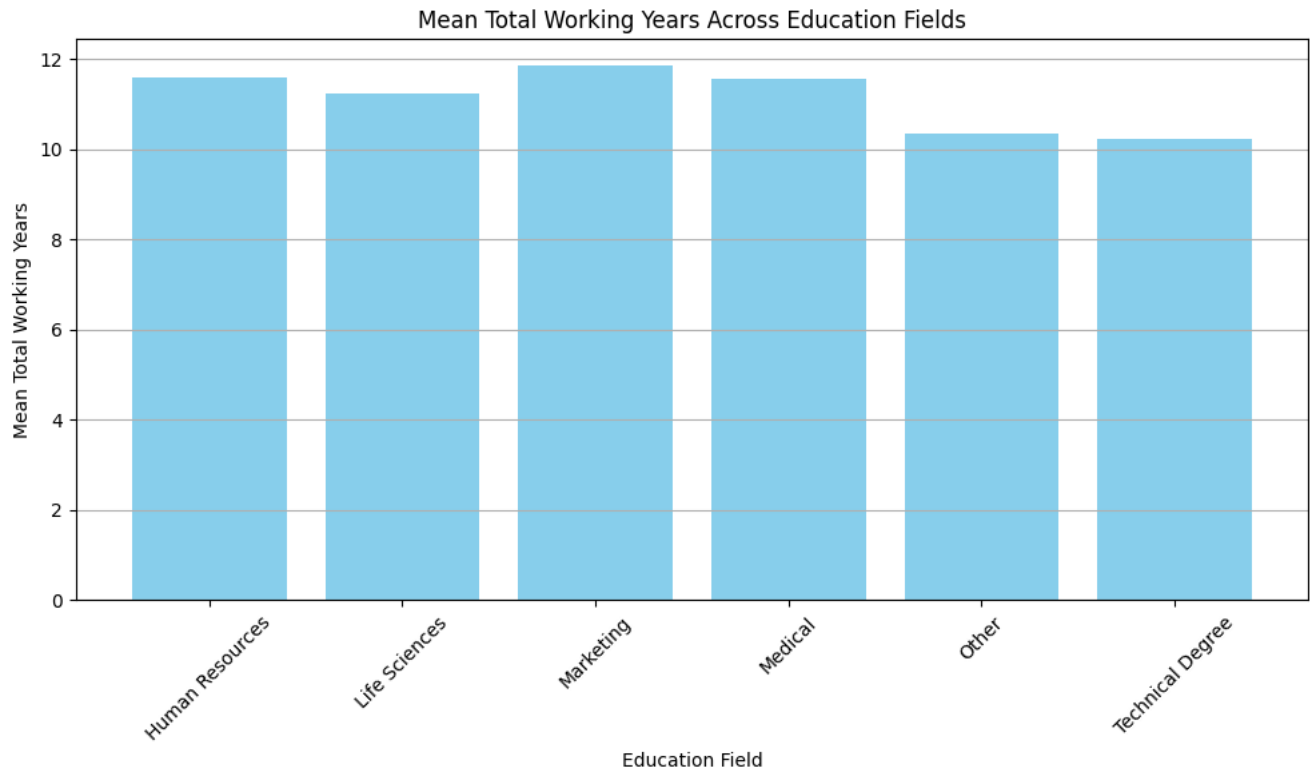
Education Field and Total Working Years

```
education_working_years = HR_data.groupby('EducationField')['TotalWorkingYears'].mean().reset_index()

plt.figure(figsize=(10, 6))
plt.bar(education_working_years['EducationField'], education_working_years['TotalWorkingYears'], color='skyblue')
plt.xlabel('Education Field')
plt.ylabel('Mean Total Working Years')
plt.title('Mean Total Working Years Across Education Fields')

# Rotating x-axis labels for better readability if needed
plt.xticks(rotation=45)

plt.grid(axis='y')
plt.tight_layout()
plt.show()
```



Insights from above graph:

- The visualization helps compare the mean total working years across different education fields.
- It allows for easy identification of which education fields tend to have longer or shorter average working experience.

Chart - 6

Pair Plot

```
sns.pairplot(HR_data)

plt.show()
```

 KeyboardInterrupt Traceback (most recent call last)

<ipython-input-15-3facd4133232> in <cell line: 1>()
 ----> 1 sns.pairplot(HR_data)

2
 3 plt.show()

⚡ 14 frames

/usr/local/lib/python3.10/dist-packages/matplotlib/cbook/_init_.py in <listcomp>(.0)

832 """Clean dead weak references from the dictionary."""

833 mapping = self._mapping

--> 834 to_drop = [key for key in mapping if key() is None]

835 for key in to_drop:

836 val = mapping.pop(key)

KeyboardInterrupt:

SEARCH STACK OVERFLOW

Error in callback <function _draw_all_if_interactive at 0x7f34c55fdcf0> (for post_execute):

 KeyboardInterrupt Traceback (most recent call last)

/usr/local/lib/python3.10/dist-packages/matplotlib/pyplot.py in _draw_all_if_interactive()

118 def _draw_all_if_interactive():

119 if matplotlib.is_interactive():

--> 120 draw_all()

121

122

⚡ 18 frames

✓ Chart - 12

Heatmap

```
correlation_data = HR_data
```

```
correlation_matrix = correlation_data.corr()
```

```
plt.figure(figsize=(20,10))
```

```
sns.heatmap(correlation_matrix,annot=True)
```

```
plt.title('Correlation Map')
```

```
plt.show()
```

<ipython-input-16-ff073efd817b>:3: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future ver
correlation_matrix = correlation_data.corr()

Correlation Map

- Several that columns have high correlation: Age, JobLevel, MonthlyIncome, NumCompaniesWorked, TotalWorkingYears, YearsAtCompany, YearsInCurrentRole, YearsSinceLastPromotion, YearsWithCurrManager

HR_data.head()

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	EducationField	EmployeeCount	EmployeeNumber
0	41	Yes	Travel_Rarely	1102	Sales		1	2	Life Sciences	1
1	49	No	Travel_Frequently	279	Research & Development		8	1	Life Sciences	1
2	37	Yes	Travel_Rarely	1373	Research & Development		2	2	Other	1
3	33	No	Travel_Frequently	1392	Research & Development		3	4	Life Sciences	1
4	27	No	Travel_Rarely	591	Research & Development		2	1	Medical	1

```
from sklearn.preprocessing import LabelEncoder
from sklearn.preprocessing import StandardScaler

HR_data_1=HR_data.copy()

le = LabelEncoder()

for i in HR_data_1.columns:
    if HR_data_1[i].dtype == 'object':
        HR_data_1[i] = le.fit_transform(HR_data_1[i])

HR_data_1.head()
```

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	EducationField	EmployeeCount	EmployeeNumber
0	41	1	2	1102	2	1	2	1	1	1
1	49	0	1	279	1	8	1	1	1	2
2	37	1	2	1373	1	2	2	4	1	4
3	33	0	1	1392	1	3	4	1	1	5
4	27	0	2	591	1	2	1	3	1	7

To get VIF, will first define x and y variables for our ML Model

```
x = HR_data_1.drop(columns=['Attrition'])
y = HR_data_1['Attrition']

from sklearn.preprocessing import StandardScaler
from statsmodels.stats.outliers_influence import variance_inflation_factor

scalar = StandardScaler()
x_scaled=scalar.fit_transform(x)

# VIF

vif = pd.DataFrame()

vif['vif']=[variance_inflation_factor(x_scaled,i) for i in range(x_scaled.shape[1])]

vif['features'] = x.columns

vif
```

return 1 - self.ssr/self.uncentered_tss

	vif	features
0	2.054226	Age
1	1.016808	BusinessTravel
2	1.026401	DailyRate
3	1.942165	Department
4	1.018096	DistanceFromHome
5	1.065295	Education
6	1.016240	EducationField
7	NaN	EmployeeCount
8	1.022699	EmployeeNumber
9	1.018022	EnvironmentSatisfaction
10	1.020038	Gender
11	1.021984	HourlyRate
12	1.020836	JobInvolvement
13	11.825039	JobLevel
14	1.894366	JobRole
15	1.022894	JobSatisfaction
16	1.844040	MaritalStatus
17	11.055038	MonthlyIncome
18	1.015803	MonthlyRate
19	1.261958	NumCompaniesWorked
20	NaN	Over18
21	1.028856	OverTime
22	2.521583	PercentSalaryHike
23	2.519973	PerformanceRating
24	1.025574	RelationshipSatisfaction
25	NaN	StandardHours
26	1.828706	StockOptionLevel
27	4.824457	TotalWorkingYears
28	1.024312	TrainingTimesLastYear
29	1.018572	WorkLifeBalance
30	4.601995	YearsAtCompany
31	2.728269	YearsInCurrentRole
32	1.678897	YearsSinceLastPromotion
33	2.782981	YearsWithCurrManager

✓ ML Model Implementation

✓ ML Model - 1

Using all Variables for ML Model-1

```
# Importing Necessary Libraries

from sklearn.preprocessing import MinMaxScaler
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from statsmodels.stats.outliers_influence import variance_inflation_factor
from sklearn.metrics import accuracy_score, confusion_matrix, roc_curve, roc_auc_score
from sklearn.metrics import classification_report
from sklearn.model_selection import GridSearchCV

HR_data_1.columns

Index(['Age', 'Attrition', 'BusinessTravel', 'DailyRate', 'Department',
      'DistanceFromHome', 'Education', 'EducationField', 'EmployeeCount',
      'EmployeeNumber', 'EnvironmentSatisfaction', 'Gender', 'HourlyRate',
      'JobInvolvement', 'JobLevel', 'JobRole', 'JobSatisfaction',
      'MaritalStatus', 'MonthlyIncome', 'MonthlyRate', 'NumCompaniesWorked',
      'Over18', 'OverTime', 'PercentSalaryHike', 'PerformanceRating',
      'RelationshipSatisfaction', 'StandardHours', 'StockOptionLevel',
      'TotalWorkingYears', 'TrainingTimesLastYear', 'WorkLifeBalance',
      'YearsAtCompany', 'YearsInCurrentRole', 'YearsSinceLastPromotion',
      'YearsWithCurrManager'],
      dtype='object')

# For ML Model 1 Using all variables from dataset, putting selected features.

x = HR_data_1[['Age', 'BusinessTravel', 'DailyRate', 'Department',
      'DistanceFromHome', 'Education', 'EducationField', 'EmployeeCount',
      'EmployeeNumber', 'EnvironmentSatisfaction', 'Gender', 'HourlyRate',
      'JobInvolvement', 'JobLevel', 'JobRole', 'JobSatisfaction',
      'MaritalStatus', 'MonthlyIncome', 'MonthlyRate', 'NumCompaniesWorked',
      'Over18', 'OverTime', 'PercentSalaryHike', 'PerformanceRating',
      'RelationshipSatisfaction', 'StandardHours', 'StockOptionLevel',
      'TotalWorkingYears', 'TrainingTimesLastYear', 'WorkLifeBalance',
      'YearsAtCompany', 'YearsInCurrentRole', 'YearsSinceLastPromotion',
      'YearsWithCurrManager']]

y = HR_data_1['Attrition']

# splitting data into train and test set.

x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.25, random_state=348)

# Transforming data standardization

scaler = MinMaxScaler()
x_train = scaler.fit_transform(x_train)
x_test = scaler.fit_transform(x_test)

# Fitting Logistic Regression model to dataset

logistic_reg = LogisticRegression()

logistic_reg.fit(x_train, y_train)

# Make predictions on the test set
y_pred = logistic_reg.predict(x_test)

# Evaluate the model using various metrics
accuracy = accuracy_score(y_test, y_pred)
confusion = confusion_matrix(y_test, y_pred)
classification_report_str = classification_report(y_test, y_pred)

# Print the evaluation metrics
print("Accuracy For ML Model 1:", accuracy)
print("Confusion Matrix For ML Model 1:\n", confusion)
print("Classification Report for ML Model 1:\n", classification_report_str)

# Training Min-Max Scaler
```

```
# Initialize MinMaxScaler
scaler = MinMaxScaler()

# Fit and transform the data using the scaler
x_train_scaled = scaler.fit_transform(x_train)
x_test_scaled = scaler.transform(x_test)

# Define the hyperparameter grid
param_grid = {
    'C': [0.001, 0.01, 0.1, 1, 10, 100], # Inverse of regularization strength
    'penalty': ['l1', 'l2'], # Regularization penalty
    'solver': ['liblinear', 'lbfgs'], # Solver for optimization
}

# Initialize GridSearchCV
grid_search = GridSearchCV(estimator=logistic_reg, param_grid=param_grid, scoring='accuracy', cv=5)

# Fit the grid search to your training data
grid_search.fit(x_train_scaled, y_train)

# Get the best hyperparameters
best_params = grid_search.best_params_

# Get the best model
best_model = grid_search.best_estimator_

# Make predictions on the test set using the best model
y_pred = best_model.predict(x_test_scaled)

# Evaluate the model with the best hyperparameters
accuracy = accuracy_score(y_test, y_pred)
confusion = confusion_matrix(y_test, y_pred)
classification_report_str = classification_report(y_test, y_pred)

# Print the best hyperparameters, best model, and evaluation metrics
print("Best Hyperparameters:", best_params)
print("Best Model:", best_model)
print("Accuracy For ML Model 1(with Hyperparameter Tuning):", accuracy)
print("Confusion Matrix For ML Model 1(with Hyperparameter Tuning):\n", confusion)
print("Classification Report For ML Model 1(with Hyperparameter Tuning):\n", classification_report_str)
```

```
Classification Report for ML Model 1:
              precision    recall  f1-score   support

     0       0.85         0.98         0.91         297
     1       0.79         0.27         0.40          71

 accuracy         0.85         368
 macro avg       0.82         0.63         0.66         368
weighted avg       0.84         0.85         0.81         368
```

```
Best Hyperparameters: {'C': 1, 'penalty': 'l1', 'solver': 'liblinear'}
Best Model: LogisticRegression(C=1, penalty='l1', solver='liblinear')
Accuracy For ML Model 1(with Hyperparameter Tuning): 0.845108695652174
Confusion Matrix For ML Model 1(with Hyperparameter Tuning):
[[293   4]
 [ 53  18]]
```

```
Classification Report for ML Model 1(with Hyperparameter Tuning):
              precision    recall  f1-score   support

     0       0.85         0.99         0.91         297
     1       0.82         0.25         0.39          71

 accuracy         0.85         368
 macro avg       0.83         0.62         0.65         368
weighted avg       0.84         0.85         0.81         368
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (stat
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_validation.py:378: FitFailedWarning:
30 fits failed out of a total of 120.
```

The score on these train-test partitions for these parameters will be set to nan.

If these failures are not expected, you can try to debug them by setting error_score='raise'.

```

/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_validation.py", line 606, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
File "/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py", line 1162, in fit
    solver = _check_solver(self.solver, self.penalty, self.dual)
File "/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py", line 54, in _check_solver
    raise ValueError(
ValueError: Solver lbfgs supports only 'l2' or 'none' penalties, got l1 penalty.

warnings.warn(some_fits_failed_message, FitFailedWarning)
/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_search.py:952: UserWarning: One or more of the test scores are nan
0.84936652 0.84936652 0.85299877      nan 0.85844097 0.8566269
0.87748252      nan 0.87658577 0.87295763 0.8747717      nan
0.87386672 0.87386672 0.87386261      nan 0.8747717 0.8747717 ]
warnings.warn(

```

Insights from ML Model 1:

- Before Tuning:

Accuracy: 0.8451

Classification Report Highlights:

Precision for label 0: 0.85

Precision for label 1: 0.79

Recall for label 0: 0.98

Recall for label 1: 0.27

F1-score for label 0: 0.91

F1-score for label 1: 0.40

- After Tuning:

Accuracy: 0.8451 (same as before tuning)

Classification Report Highlights:

Precision for label 0: 0.85

Precision for label 1: 0.82

Recall for label 0: 0.99

Recall for label 1: 0.25

F1-score for label 0: 0.91

F1-score for label 1: 0.39

The accuracy of the model remained the same, indicating that the tuning did not significantly impact overall correct predictions. Precision Improved for label 1 after tuning, which means that among the predicted positives, a higher proportion were actually true positives.

✓ ML Model - 2

Using selected Variables for ML Model-2 (Neglecting verible whose VIF is more than 5)

```

# For ML Model 1 Using all variables from dataset, putting selected features.

x = HR_data_1[['Age', 'BusinessTravel', 'DailyRate', 'Department',
              'DistanceFromHome', 'Education', 'EducationField', 'EmployeeCount',
              'EmployeeNumber', 'EnvironmentSatisfaction', 'Gender', 'HourlyRate',
              'JobInvolvement', 'JobRole', 'JobSatisfaction',
              'MaritalStatus', 'MonthlyRate', 'NumCompaniesWorked',
              'Over18', 'OverTime', 'PercentSalaryHike', 'PerformanceRating',
              'RelationshipSatisfaction', 'StandardHours', 'StockOptionLevel',
              'TrainingTimesLastYear', 'WorkLifeBalance',
              'YearsInCurrentRole', 'YearsSinceLastPromotion',
              'YearsWithCurrManager']]

y = HR_data_1['Attrition']

# splitting data into train and test set.

x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.25,random_state=348)

# Transforming data standardization

scaler = MinMaxScaler()
x_train = scaler.fit_transform(x_train)

```

```

x_test = scaler.fit_transform(x_test)

# Fitting Logistic Regression model to dataset

logistic_reg = LogisticRegression()

logistic_reg.fit(x_train,y_train)

# Make predictions on the test set
y_pred = logistic_reg.predict(x_test)

# Evaluate the model using various metrics
accuracy = accuracy_score(y_test, y_pred)
confusion = confusion_matrix(y_test, y_pred)
classification_report_str = classification_report(y_test, y_pred)

# Print the evaluation metrics
print("Accuracy For ML Model 1:", accuracy)
print("Confusion Matrix For ML Model 1:\n", confusion)
print("Classification Report for ML Model 1:\n", classification_report_str)

# Initialize MinMaxScaler
scaler = MinMaxScaler()

# Fit and transform the data using the scaler
x_train_scaled = scaler.fit_transform(x_train)
x_test_scaled = scaler.transform(x_test)

# Define the hyperparameter grid
param_grid = {
    'C': [0.001, 0.01, 0.1, 1, 10, 100], # Inverse of regularization strength
    'penalty': ['l1', 'l2'], # Regularization penalty
    'solver': ['liblinear', 'lbfgs'], # Solver for optimization
}

# Initialize GridSearchCV
grid_search = GridSearchCV(estimator=logistic_reg, param_grid=param_grid, scoring='accuracy', cv=5)

# Fit the grid search to your training data
grid_search.fit(x_train_scaled, y_train)

# Get the best hyperparameters
best_params = grid_search.best_params_

# Get the best model
best_model = grid_search.best_estimator_

# Make predictions on the test set using the best model
y_pred = best_model.predict(x_test_scaled)

# Evaluate the model with the best hyperparameters
accuracy = accuracy_score(y_test, y_pred)
confusion = confusion_matrix(y_test, y_pred)
classification_report_str = classification_report(y_test, y_pred)

# Print the best hyperparameters, best model, and evaluation metrics
print("Best Hyperparameters:", best_params)
print("Best Model:", best_model)
print("Accuracy For ML Model 1(with Hyperparameter Tuning):", accuracy)
print("Confusion Matrix For ML Model 1(with Hyperparameter Tuning):\n", confusion)
print("Classification Report for ML Model 1(with Hyperparameter Tuning):\n", classification_report_str)

Accuracy For ML Model 1: 0.842391304347826
Confusion Matrix For ML Model 1:
[[290  7]
 [ 51 20]]
Classification Report for ML Model 1:

```

	precision	recall	f1-score	support
0	0.85	0.98	0.91	297
1	0.74	0.28	0.41	71
accuracy			0.84	368
macro avg	0.80	0.63	0.66	368
weighted avg	0.83	0.84	0.81	368

```

Best Hyperparameters: {'C': 1, 'penalty': 'l2', 'solver': 'liblinear'}
Best Model: LogisticRegression(C=1, solver='liblinear')
Accuracy For ML Model 1(with Hyperparameter Tuning): 0.842391304347826
Confusion Matrix For ML Model 1(with Hyperparameter Tuning):
[[290  7]

```

```
[ 51 20]]
Classification Report for ML Model 1(with Hyperparameter Tuning):
      precision    recall  f1-score   support

      0       0.85       0.98       0.91       297
      1       0.74       0.28       0.41        71

 accuracy         0.84       368
 macro avg       0.80       0.63       0.66       368
 weighted avg    0.83       0.84       0.81       368

/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_validation.py:378: FitFailedWarning:
30 fits failed out of a total of 120.
The score on these train-test partitions for these parameters will be set to nan.
If these failures are not expected, you can try to debug them by setting error_score='raise'.

Below are more details about the failures:
-----
30 fits failed with the following error:
Traceback (most recent call last):
  File "/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_validation.py", line 686, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
  File "/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py", line 1162, in fit
    solver = _check_solver(self.solver, self.penalty, self.dual)
  File "/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py", line 54, in _check_solver
    raise ValueError(
ValueError: Solver lbfgs supports only 'l2' or 'none' penalties, got l1 penalty.

    warnings.warn(some_fits_failed_message, FitFailedWarning)
/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_search.py:952: UserWarning: One or more of the test scores are non
0.84936652 0.84936652 0.85299877      nan 0.85844097 0.85572193
0.86569724      nan 0.86751131 0.86751131 0.86661045      nan
0.86569724 0.86478815 0.86661045      nan 0.86661045 0.86661045]
    warnings.warn(
```

Insights from ML Model 2

- Before Tuning:
Accuracy: 0.842 (84.2%)
- After Tuning:
Accuracy: 0.842 (84.2%)

The performance metrics (accuracy, confusion matrix, and classification report) remain unchanged before and after hyperparameter tuning.

ML Model - 3

Decision Tree Model




```
# Importing decision tree classifier

from sklearn.tree import DecisionTreeClassifier

# For ML Model 3, using selected variable from previous model which gives more accuracy.

x = HR_data_1[['Age', 'BusinessTravel', 'DailyRate', 'Department',
               'DistanceFromHome', 'Education', 'EducationField', 'EmployeeCount',
               'EmployeeNumber', 'EnvironmentSatisfaction', 'Gender', 'HourlyRate',
               'JobInvolvement', 'JobRole', 'JobSatisfaction',
               'MaritalStatus', 'MonthlyRate', 'NumCompaniesWorked',
               'Over18', 'OverTime', 'PercentSalaryHike', 'PerformanceRating',
               'RelationshipSatisfaction', 'StandardHours', 'StockOptionLevel',
               'TrainingTimesLastYear', 'WorkLifeBalance',
               'YearsInCurrentRole', 'YearsSinceLastPromotion',
               'YearsWithCurrManager']]

y = HR_data_1['Attrition']

# splitting data into train and test set.

x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.25,random_state=348)

# Transforming data standardization

scaler = MinMaxScaler()
x_train = scaler.fit_transform(x_train)
x_test = scaler.fit_transform(x_test)

# Fitting Logistic Regression model to dataset

decision_tree = DecisionTreeClassifier()
decision_tree.fit(x_train, y_train)

# Make predictions on the test set
y_pred = decision_tree.predict(x_test)

# Evaluate the model using various metrics
accuracy = accuracy_score(y_test, y_pred)
confusion = confusion_matrix(y_test, y_pred)
classification_report_str = classification_report(y_test, y_pred)

# Print the evaluation metrics
print("Accuracy For Decision Tree Model:", accuracy)
print("Confusion Matrix Decision Tree Model:\n", confusion)
print("Classification Report Decision Tree Model:\n", classification_report_str)

# Initialize MinMaxScaler
scaler = MinMaxScaler()

# Fit and transform the data using the scaler
x_train_scaled = scaler.fit_transform(x_train)
x_test_scaled = scaler.transform(x_test)

# Define the hyperparameter grid

param_grid = {
    'criterion': ['gini', 'entropy'],
    'max_depth': range(10,15),
    'min_samples_leaf': range(2,6),
    'min_samples_split': range(3,8),
    'max_leaf_nodes': range(5,10)
}

# Initialize GridSearchCV
grid_search = GridSearchCV(estimator=decision_tree, param_grid=param_grid, scoring='accuracy', cv=5)

# Fit the grid search to your training data
grid_search.fit(x_train_scaled, y_train)

# Get the best hyperparameters
best_params = grid_search.best_params_

# Get the best model
best_model = grid_search.best_estimator_

# Make predictions on the test set using the best model
```

```

y_pred = best_model.predict(x_test_scaled)

# Evaluate the model with the best hyperparameters
accuracy = accuracy_score(y_test, y_pred)
confusion = confusion_matrix(y_test, y_pred)
classification_report_str = classification_report(y_test, y_pred)

# Print the best hyperparameters, best model, and evaluation metrics
print("Best Hyperparameters:", best_params)
print("Best Model:", best_model)
print("Accuracy For Decision Tree Model (with Hyperparameter Tuning):", accuracy)
print("Confusion Matrix For Decision Tree Model (with Hyperparameter Tuning):\n", confusion)
print("Classification Report for Decision Tree Model (with Hyperparameter Tuning):\n", classification_report_str)

Accuracy For Decision Tree Model: 0.7608695652173914
Confusion Matrix Decision Tree Model:
[[251  46]
 [ 42  29]]
Classification Report Decision Tree Model:

```

	precision	recall	f1-score	support
0	0.86	0.85	0.85	297
1	0.39	0.41	0.40	71
accuracy			0.76	368
macro avg	0.62	0.63	0.62	368
weighted avg	0.77	0.76	0.76	368

```

Best Hyperparameters: {'criterion': 'gini', 'max_depth': 10, 'max_leaf_nodes': 6, 'min_samples_leaf': 2, 'min_samples_split': 3}
Best Model: DecisionTreeClassifier(max_depth=10, max_leaf_nodes=6, min_samples_leaf=2,
min_samples_split=3)
Accuracy For Decision Tree Model (with Hyperparameter Tuning): 0.8260869565217391
Confusion Matrix For Decision Tree Model (with Hyperparameter Tuning):
[[289  8]
 [ 56 15]]
Classification Report for Decision Tree Model (with Hyperparameter Tuning):

```

	precision	recall	f1-score	support
0	0.84	0.97	0.90	297
1	0.65	0.21	0.32	71
accuracy			0.83	368
macro avg	0.74	0.59	0.61	368
weighted avg	0.80	0.83	0.79	368

Insights from Decision Tree Model-

- Before Tuning:

Accuracy: 76.08%

Classification Report:

Precision for class 0: 86%

Recall for class 0: 85%

F1-score for class 0: 85%

Precision for class 1: 39%

Recall for class 1: 41%

F1-score for class 1: 40%

- After Tuning:

Accuracy: 82.61%

Classification Report:

Precision for class 0: 84%

Recall for class 0: 97%

F1-score for class 0: 90%

Precision for class 1: 65%

Recall for class 1: 21%

F1-score for class 1: 32%

The model's accuracy increased from 76.08% to 82.61% after hyperparameter tuning. This indicates a better overall prediction capability. The model after tuning performs better in terms of overall accuracy and correctly identifying class 0 instances (higher precision, recall, and F1-

score).

✓ ML Model - 4

RandomForestClassifier

```
# Importing decision tree classifier

from sklearn.tree import DecisionTreeClassifier

# For ML Model RandomForestClassifier, using selected variable from previous model which gives more accuracy.

x = HR_data_1[['Age', 'BusinessTravel', 'DailyRate', 'Department',
               'DistanceFromHome', 'Education', 'EducationField', 'EmployeeCount',
               'EmployeeNumber', 'EnvironmentSatisfaction', 'Gender', 'HourlyRate',
               'JobInvolvement', 'JobRole', 'JobSatisfaction',
               'MaritalStatus', 'MonthlyRate', 'NumCompaniesWorked',
               'Over18', 'OverTime', 'PercentSalaryHike', 'PerformanceRating',
               'RelationshipSatisfaction', 'StandardHours', 'StockOptionLevel',
               'TrainingTimesLastYear', 'WorkLifeBalance',
               'YearsInCurrentRole', 'YearsSinceLastPromotion',
               'YearsWithCurrManager']]

y = HR_data_1['Attrition']

# splitting data into train and test set.

x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.25,random_state=348)

# Transforming data standardization

scaler = MinMaxScaler()
x_train = scaler.fit_transform(x_train)
x_test = scaler.fit_transform(x_test)

# Fitting RandomForest model, first will import from sklearn package

from sklearn.ensemble import RandomForestClassifier

random_for = RandomForestClassifier()

random_for.fit(x_train,y_train)

# Make predictions on the test set
y_pred = random_for.predict(x_test)

# Evaluate the model using various metrics
accuracy = accuracy_score(y_test, y_pred)
confusion = confusion_matrix(y_test, y_pred)
classification_report_str = classification_report(y_test, y_pred)

# Print the evaluation metrics
print("Accuracy For RandomForestClassifier:", accuracy)
print("Confusion Matrix RandomForestClassifier:\n", confusion)
print("Classification Report RandomForestClassifier:\n", classification_report_str)

# Initialize MinMaxScaler
scaler = MinMaxScaler()

# Fit and transform the data using the scaler
x_train_scaled = scaler.fit_transform(x_train)
x_test_scaled = scaler.transform(x_test)

# Define the hyperparameter grid

param_grid = {
    'criterion': ['gini'],
    'max_depth': range(10,12),
    'min_samples_leaf': range(2,4),
    'min_samples_split': range(3,8),
    'max_leaf_nodes': range(5,10)
}

# Initialize GridSearchCV
grid_search = GridSearchCV(estimator=random_for, param_grid=param_grid, scoring='accuracy', cv=5)

# Fit the grid search to your training data
grid_search.fit(x_train_scaled, y_train)

# Get the best hyperparameters
best_params = grid_search.best_params_

# Get the best model
best_model = grid_search.best_estimator_
```

```
# Make predictions on the test set using the best model
y_pred = best_model.predict(x_test_scaled)

# Evaluate the model with the best hyperparameters
accuracy = accuracy_score(y_test, y_pred)
confusion = confusion_matrix(y_test, y_pred)
classification_report_str = classification_report(y_test, y_pred)

# Print the best hyperparameters, best model, and evaluation metrics
print("Best Hyperparameters:", best_params)
print("Best Model:", best_model)
print("Accuracy For RandomForestClassifier (with Hyperparameter Tuning):", accuracy)
print("Confusion Matrix For RandomForestClassifier (with Hyperparameter Tuning):\n", confusion)
print("Classification Report for RandomForestClassifier (with Hyperparameter Tuning):\n", classification_report_str)

Accuracy For RandomForestClassifier: 0.8315217391304348
Confusion Matrix RandomForestClassifier:
[[295  2]
 [ 60 11]]
Classification Report RandomForestClassifier:
      precision    recall  f1-score   support

     0       0.83       0.99       0.90         297
     1       0.85       0.15       0.26          71

   accuracy          0.83         368
  macro avg       0.84       0.57       0.58         368
 weighted avg       0.83       0.83       0.78         368

Best Hyperparameters: {'criterion': 'gini', 'max_depth': 11, 'max_leaf_nodes': 9, 'min_samples_leaf': 3, 'min_samples_split': 6}
Best Model: RandomForestClassifier(max_depth=11, max_leaf_nodes=9, min_samples_leaf=3,
                                min_samples_split=6)
Accuracy For RandomForestClassifier (with Hyperparameter Tuning): 0.8097826086956522
Confusion Matrix For RandomForestClassifier (with Hyperparameter Tuning):
[[297  0]
 [ 70 11]]
Classification Report for RandomForestClassifier (with Hyperparameter Tuning):
      precision    recall  f1-score   support

     0       0.81       1.00       0.89         297
     1       1.00       0.01       0.03          71

   accuracy          0.81         368
  macro avg       0.90       0.51       0.46         368
 weighted avg       0.85       0.81       0.73         368
```

Insights from ML Model 4-

- Before Tuning:
Accuracy: 83.15%
- After Tuning:
Accuracy: 80.98%

Accuracy: After tuning, the accuracy decreased slightly from 83.15% to 80.98%.

Confusion Matrix: Before tuning, the model misclassified 62 instances. After tuning, it misclassified 70 instances of class 1.

Precision, Recall, F1-score: For class 0 (majority class), precision and recall remained fairly similar after tuning. For class 1, precision was perfect (1.00) after tuning but recall dropped significantly from 0.15 to almost 0.01, indicating the model's inability to correctly identify instances of class 1.

Before Tuning: The model had a higher overall accuracy but struggled with recall and precision for the minority class (class 1). After Tuning: Although the accuracy decreased slightly, there was an improvement in precision for class 1 but at the cost of significantly lower recall.

✓ ML Model - 5

KNeighborsClassifier

```
# Importing decision tree classifier

from sklearn.tree import DecisionTreeClassifier

# For ML Model 7, Using all KNeighborsClassifier

x = HR_data_1[['Age', 'BusinessTravel', 'DailyRate', 'Department',
               'DistanceFromHome', 'Education', 'EducationField', 'EmployeeCount',
               'EmployeeNumber', 'EnvironmentSatisfaction', 'Gender', 'HourlyRate',
               'JobInvolvement', 'JobRole', 'JobSatisfaction',
               'MaritalStatus', 'MonthlyRate', 'NumCompaniesWorked',
               'Over18', 'OverTime', 'PercentSalaryHike', 'PerformanceRating',
               'RelationshipSatisfaction', 'StandardHours', 'StockOptionLevel',
               'TrainingTimesLastYear', 'WorkLifeBalance',
               'YearsInCurrentRole', 'YearsSinceLastPromotion',
               'YearsWithCurrManager']]

y = HR_data_1['Attrition']

# splitting data into train and test set.

x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.25, random_state=348)

# Transforming data standardization

scaler = MinMaxScaler()
x_train = scaler.fit_transform(x_train)
x_test = scaler.fit_transform(x_test)

# Fitting KNeighborsClassifier model, first will import from sklearn package

from sklearn.neighbors import KNeighborsClassifier

kNN = KNeighborsClassifier()

kNN.fit(x_train, y_train)

# Make predictions on the test set
y_pred = kNN.predict(x_test)

# Evaluate the model using various metrics
accuracy = accuracy_score(y_test, y_pred)
confusion = confusion_matrix(y_test, y_pred)
classification_report_str = classification_report(y_test, y_pred)

# Print the evaluation metrics
print("Accuracy For KNeighborsClassifier:", accuracy)
print("Confusion Matrix for KNeighborsClassifier:\n", confusion)
print("Classification Report for KNeighborsClassifier:\n", classification_report_str)

# Initialize MinMaxScaler
scaler = MinMaxScaler()

# Fit and transform the data using the scaler
```